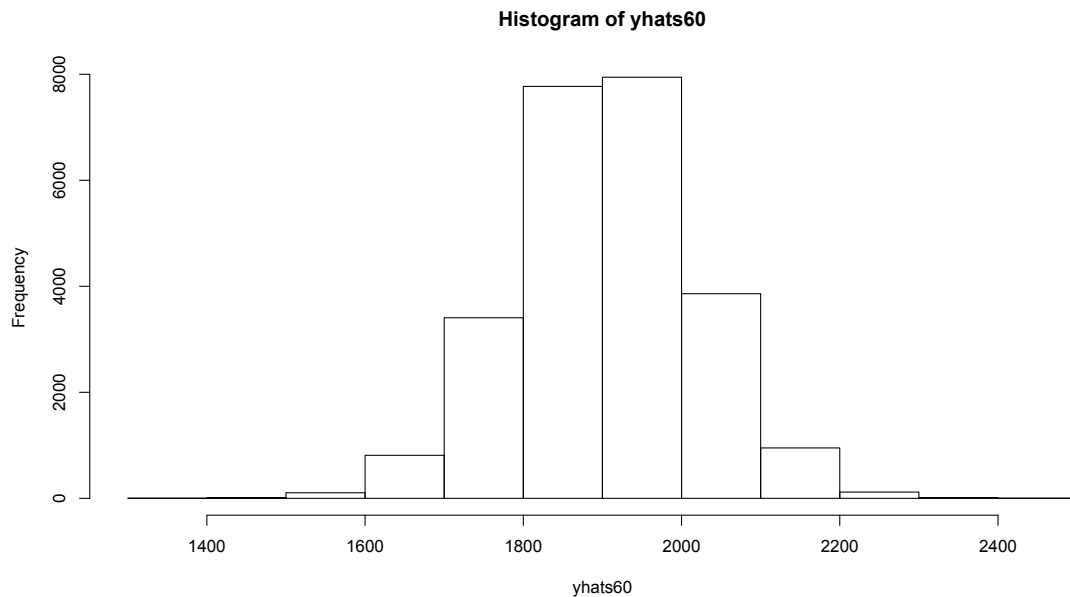


COMMENTS FROM GRADED EXAM: The 95% credible interval for the posterior predictive distribution of IgG measurements for the next person who enters with uptake $x=60$ was too narrow. The error was that I took the credible interval of the N regression means instead drawing N random normal values with each mean as a center and each corresponding sigsq value as the variance.



Mean: 1904.294

95% Credible Interval: (1678.518, 2128.151)

#Read file and attach

```
data = read.csv(file.choose(), header=TRUE)
```

```
attach(data)
```

#Set gibbs parameters

```
y = data$IgG
```

```
x = data$Max.O2.Uptake
```

```
output = summary(glm(y ~ x + I(x^2)))
```

```
N = 25000
```

```
lag = 500
```

```
burnin = 0
```

```
b0 = output$coef[1,1]
```

```
b1 = output$coef[2,1]
```

```
b2 = output$coef[3,1]
```

#Gibbs function receives eight parameters: "x" and "y" data values; "b0", "b1", and "b2" prior regression coefficients, "N" for number of realizations, "lag" for determining how many realizations to skip between saves, and "burnin" for determining how many realizations to skip before starting to save. Gibbs generates N independent random normal values for each regression coefficient and N independent inverse gamma values for the regression variance.

```
gibbs <- function(x,y,b0,b1,b2,N,lag,burnin) {

  #obtain length of data
  n = length(x)
  #Set N to be N*lag+burnin
  N <- N*lag + burnin

  #Initialize vectors to hold the beta coefficients and sigsq realizations
  b0s = NULL
  b1s = NULL
  b2s = NULL
  s2s = NULL

  for(i in 1:N) {

    #Generate a sigsq, s2, based on current regression coefficients b0, b1,
    b2
    s2 = 1/rgamma(1, n/2, sum((y-b0-b1*x-b2*x^2)^2)/2)
    b0 = rnorm(1, sum(y-b1*x-b2*x^2)/n, sqrt(s2/n))
    b1 = rnorm(1, sum(x*(y-b0-b2*x^2))/sum(x^2), sqrt(s2/sum(x^2)))
    b2 = rnorm(1, sum(x^2*(y-b0-b1*x))/sum(x^4), sqrt(s2/sum(x^4)))

    #if i is greater than burnin and if i is a multiple of the lag, store alpha,
    beta, and sigsq
    if(i > burnin) {
      if(i %% lag == 0) {
        b0s <- c(b0s,b0)
        b1s <- c(b1s,b1)
        b2s <- c(b2s,b2)
        s2s <- c(s2s,s2)
      }
    }

    }

  vectors <- list("beta0s" = b0s, "beta1s" = b1s, "beta2s" = b2s, "sigsqs" = s2s)
  return(vectors)
}
```

#ii USE GIBBS TO SAMPLE FROM THE MARGINAL POSTERIOR OF EACH OF THE FOUR PARAMETERS

```
v = gibbs(x,y,b0,b1,b2,N,lag,burnin)
```

```
beta0s = v$beta0s
```

```
beta1s = v$beta1s
```

```
beta2s = v$beta2s
```

```
sigsqs = v$sigsqs
```

```
mean(beta0s)
```

```
#[1] -1452.651
```

```
mean(beta1s)
```

```
#[1] 87.85156
```

```
mean(beta2s)
```

```
#[1] -0.5320706
```

```
mean(sigsqs)
```

```
#[1] 12219.16
```

```
#Plot each parameter with ts.plot, hist, and acf
```

```
par(mfrow=c(3,1)) #split plotting window into 3 rows and 1 column
```

```
ts.plot(beta0s,xlab="Iterations")
```

```
hist(beta0s,probability=T, cex.lab=1.5, cex.axis=1.5)
```

```
acf(beta0s,lag.max=500)
```

```
par(mfrow=c(3,1)) #split plotting window into 3 rows and 1 column
```

```
ts.plot(beta1s,xlab="Iterations")
```

```
hist(beta1s,probability=T, cex.lab=1.5, cex.axis=1.5)
```

```
acf(beta1s,lag.max=500)
```

```
par(mfrow=c(3,1)) #split plotting window into 3 rows and 1 column
```

```
ts.plot(beta2s,xlab="Iterations")
```

```
hist(beta2s,probability=T, cex.lab=1.5, cex.axis=1.5)
```

```
acf(beta2s,lag.max=500)
```

```
par(mfrow=c(3,1)) #split plotting window into 3 rows and 1 column
```

```
ts.plot(sigsqs,xlab="Iterations")
```

```
hist(sigsqs,probability=T, cex.lab=1.5, cex.axis=1.5)
```

```
acf(sigsqs,lag.max=500)
```

```
#Create a matrix of regression coefficients - N rows and 3 columns - using the posterior distribution of each beta (USE THE POSTERIOR TO SET THE CURVES)
```

```
k=1
```

```
mcoef <- matrix(, nrow = N, ncol = 3)
```

```
for(i in 1:N) {
```

```

        coord = c(beta0s[i], beta1s[i], beta2s[i])
        mcoef[k,] = coord
        k = k + 1
    }

#iii CREATE N CURVES AND SELECT THREE AT RANDOM TO SUPERIMPOSE THREE
QUADRATIC CURVES
#Create a matrix of yhat values - N rows and 30 columns (30 x values) using the
matrix of regression coefficients
yhats <- matrix(, nrow = N, ncol = length(x))
for(i in 1:N) {
    yhat <- mcoef[i,1] + mcoef[i,2]*x + mcoef[i,3]*x^2
    yhats[i,] = yhat
}

#Select 3 random numbers between 0 and 25000 to represent the indices of the 3
randomly selected coefficient coordinates from the mcoef matrix
rLines = sample(1:25000, 3, replace=F)

#Create a 3-row matrix of yhat values for each x using the randomly selected
coordinates
yhats3 <- matrix(, nrow = 3, ncol = length(x))
for(i in 1:3) {
    yhat <- mcoef[rLines[i],1] + mcoef[rLines[i],2]*x + mcoef[rLines[i],3]*x^2
    yhats3[i,] = yhat
}

#iv SUPERIMPOSE POINTWISE MEAN REGRESSION CURVE
#Calculate the means of the columns (2) of the yhats matrix
means = apply(yhats, 2, mean)

#Plot x and y with possible fitted curves after spline smoothing
smooth1 = smooth.spline(x,yhats3[1,], spar=0.35)
smooth2 = smooth.spline(x,yhats3[2,], spar=0.35)
smooth3 = smooth.spline(x,yhats3[3,], spar=0.35)
smooth4 = smooth.spline(x, means, spar=0.35)

plot(x,y)
lines(smooth1, lty=2)
lines(smooth2, lty=2)
lines(smooth3, lty=2)
lines(smooth4, lwd=2.5)

#DRAW N REALIZATIONS FROM THE POSTERIOR PREDICTIVE DISTRIBUTION OF
IgG MEASUREMENTS FOR THE NEXT PERSON WHO ENTERS WITH UPTAKE x=60

```

```
#Using the matrix of regression coefficients - mcoef - create a vector of mean yhat values with O2 uptake = 60 - each yhat value is the mean of a normal distribution (USE THE CURVE TO SET THE MEANS FOR EACH CURVE)
```

```
ymeans60 = NULL
```

```
for(i in 1:N) {
```

```
  yhat60 <- mcoef[i,1] + mcoef[i,2]*60 + mcoef[i,3]*60^2
```

```
  ymeans60 = c(ymeans60, yhat60)
```

```
}
```

```
#Using each mean yhat value as the center of a normal distribution, draw draw from each one random normal value N realizations - one for each curve (USE EACH MEAN TO GENERATE A SINGLE REALIZATION YIELDING N TOTAL REALIZATIONS)
```

```
yhats60 = NULL
```

```
for(i in 1:N) {
```

```
  yhats60 <- c(yhats60, rnorm(1, ymeans60[i], sqrt(sigsqs[i])))
```

```
}
```

```
hist(yhats60)
```

```
mean(yhats60)
```

```
#[1] 1904.294
```

```
quantile(yhats60, 0.025)
```

```
#[1] 1678.518
```

```
quantile(yhats60, 0.975)
```

```
#[1] 2128.151
```

```
detach(data)
```